

The goal of this project is to build image classification model based on door numbers from Street View images. We will first use KNN and then use Neural networks and compare the efficiency and accuracy of the traditional model and neural networks.

The model should be able to take an image from the SVHN database as identify the digit. This is a multi-class classification problem with 10 classes - 0-9 digits.

In [28]:

```
import numpy as np
from sklearn.neighbors import KNeighborsClassifier
import pandas as pd
import h5py
from sklearn.model_selection import train_test_split
import cv2
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
```

In [29]:

```
dataFromFile = h5py.File('SVHN_single_grey1.h5', 'r')

X_train = dataFromFile['X_train']
X_train_original = dataFromFile['X_train']
X_test = dataFromFile['X_test']
y_test = dataFromFile['y_test']
y_train = dataFromFile['y_train']

X_train = X_train[:].reshape(-1, 1024) / 255
X_test = X_test[:].reshape(-1, 1024) / 255
```

In [30]:

```
# visualizing the first 10 images in the dataset and their labels
%matplotlib inline
import matplotlib.pyplot as plt
plt.figure(figsize=(10, 1))
for i in range(10):
    plt.subplot(1, 10, i+1)
    plt.imshow(X_train_original[i], cmap="gray")
    plt.axis('off')
plt.show()
print('label for each of the above image: %s' % (y_train[0:10]))
```



label for each of the above image: [2 6 7 4 4 0 3 0 7 3]

We will now consider a subset of the data for KNN classification. K Nearest neighbours classification is a traditional ML model building. We will analyze it is pertinence in the context of OCR.

In [31]:

```
X_train_sub = X_train[:3000]
y_train_sub = y_train[:3000]
X_test_sub = X_test[:1000]
y_test_sub = y_test[:1000]
```

In [32]:

```
for k in range(1,30):
    knnClassifier = KNeighborsClassifier(n_neighbors=k, weights = 'uniform', met
ric='euclidean')
    knnClassifier.fit(X_train_sub, y_train_sub)

    score = knnClassifier.score(X_test_sub, y_test_sub)
    print(f"When k={k} score is {score*100}")
```

```
When k=1 score is 32.1
When k=2 score is 30.7
When k=3 score is 30.7
When k=4 score is 33.6
When k=5 score is 33.300000000000004
When k=6 score is 33.6
When k=7 score is 32.800000000000004
When k=8 score is 34.4
When k=9 score is 34.599999999999994
When k=10 score is 34.300000000000004
When k=11 score is 34.5
When k=12 score is 34.599999999999994
When k=13 score is 34.4
When k=14 score is 35.4
When k=15 score is 34.699999999999996
When k=16 score is 35.8
When k=17 score is 35.4
When k=18 score is 35.5
When k=19 score is 35.0
When k=20 score is 34.8
When k=21 score is 34.8
When k=22 score is 34.0
When k=23 score is 33.300000000000004
When k=24 score is 34.1
When k=25 score is 33.5
When k=26 score is 33.2
When k=27 score is 33.6
When k=28 score is 33.2
When k=29 score is 33.2
```

The accuracy seems to be the best at 35.8% for k value 16. Let's set the k value to 16 and train a model with the entire training dataset.

In [33]:

```
knnClassifier = KNeighborsClassifier(n_neighbors=16)
knnClassifier.fit(X_train, y_train)

predictions_knn = knnClassifier.predict(X_test)
```

In [34]:

```
print("EVALUATION ON TESTING DATA")
print(classification_report(y_test, predictions_knn))

print("Confusion matrix")
print(confusion_matrix(y_test, predictions_knn))
```

EVALUATION ON TESTING DATA

	precision	recall	f1-score	support
0	0.44	0.69	0.54	1814
1	0.47	0.73	0.57	1828
2	0.61	0.55	0.58	1803
3	0.45	0.43	0.44	1719
4	0.66	0.65	0.65	1812
5	0.51	0.39	0.44	1768
6	0.49	0.41	0.44	1832
7	0.71	0.62	0.66	1808
8	0.46	0.36	0.40	1812
9	0.56	0.42	0.48	1804
accuracy			0.53	18000
macro avg	0.54	0.52	0.52	18000
weighted avg	0.54	0.53	0.52	18000

Confusion matrix

```
[[1254  69  35  39  53  46 103  35  81  99]
 [  97 1336  65  88  79  35  29  48  20  31]
 [ 101  230 987  98  46  34  37 141  52  77]
 [ 132  245 143 739  50 156  35  58 101  60]
 [ 108  240  43  59 1180  15  57  19  39  52]
 [ 166  164  59 258  62 686 152  31 121  69]
 [ 323  123  36  67 130 121 742  25 209  56]
 [  97  213 115  91  29  34  38 1123  26  42]
 [ 255  113  64 121  97 116 259  27 654 106]
 [ 335  134  66  91  68 100  61  70 119 760]]
```

As you can see above, using KNN has very low accuracy and the process with all 42000 images in the training data and 18000 images in test data took 30+ minutes. This is one of the obvious reasons the traditional classification models are not used for OCR.

Neural Network Model

We will now import the appropriate modules for developing a neural network model. We will add layers as needed and do BatchNormalization and dropouts as needed to get the best accuracy. We will use ReLU as the activation function.

In [35]:

```
import tensorflow.keras
from tensorflow.keras import losses
from tensorflow.keras import optimizers
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dropout, ReLU, Dense, BatchNormalization
```

The number of classes in the case of OCR of house numbers is 10, which includes all digits from 0-9. We will use `sparse_categorical_crossentropy` for losses, to treat the outcome as categorical values and not continuous. We will also use `'sparse_categorical_accuracy'` for assessing the metrics of the model. Since the output of the model is categorical, we will use softmax activation function for the final layer and ReLU for the hidden layers to get better accuracy.

In [88]:

```
model = Sequential()
model.add(BatchNormalization(input_shape=(1024,)))
model.add(Dense(256,activation='relu'))
model.add(BatchNormalization(input_shape=(256,)))
model.add(Dense(64,activation='relu'))
model.add(Dropout(rate=0.2))
num_classes = len(np.unique(y_train))
model.add(Dense(num_classes,activation='softmax'))
adam = optimizers.Adam(lr=0.002)
model.compile(loss=losses.sparse_categorical_crossentropy, optimizer=adam, metrics=[ 'sparse_categorical_accuracy'])
model.fit(X_train, y_train, batch_size=200, epochs=30,validation_data=(X_test, y_test),verbose=0)

_, train_data_accuracy = model.evaluate(X_train, y_train)
print('Accuracy: %.2f' % (train_data_accuracy*100))

_, test_data_accuracy = model.evaluate(X_test, y_test)
print('Accuracy: %.2f' % (test_data_accuracy*100))

42000/42000 [=====] - 4s 90us/sample - los
s: 0.2503 - sparse_categorical_accuracy: 0.9219
Accuracy: 92.19
18000/18000 [=====] - 1s 75us/sample - los
s: 0.5037 - sparse_categorical_accuracy: 0.8601
Accuracy: 86.01
```


As observed above the accuracy of the model with the training data is 92.05% and the accuracy with test data is 85.07%.

In [57]:

```
nnmodelPredictions = model.predict_classes(X_train)
```

In [63]:

```
plt.figure(figsize=(10, 1))
for i in range(10):
    plt.subplot(1, 10, i+1)
    plt.imshow(X_train_original[i], cmap="gray")
    plt.axis('off')
plt.show()
print('label for each of the above image: %s' % (y_train[0:10]))
print('Model Prediction for each of the above image: %s' % nnmodelPredictions[0:10])
```



label for each of the above image: [2 6 7 4 4 0 3 0 7 3]

Model Prediction for each of the above image: [2 5 7 4 4 9 3 7 7 3]

We will fine tune some of the hyperparameters and develop another model to see if the performance can be improved.

In [82]:

```
model2 = Sequential()
model2.add(BatchNormalization(input_shape=(1024,)))
model2.add(Dense(256,activation='relu'))
model2.add(BatchNormalization(input_shape=(256,)))
model2.add(Dense(128,activation='relu'))
model2.add(BatchNormalization(input_shape=(128,)))
model2.add(Dense(64,activation='relu'))
model2.add(Dropout(rate=0.2))
num_classes = len(np.unique(y_train))
model2.add(Dense(num_classes,activation='softmax'))
lrArr = [0.01,0.02,0.03,0.04,0.05]
momArr = [0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9]
for lr in lrArr:
    for mom in momArr:
        sgd = optimizers.SGD(lr=lr, momentum=mom, nesterov=False)
        model2.compile(loss=losses.sparse_categorical_crossentropy, optimizer=sgd, metrics=['sparse_categorical_accuracy'])

        model2.fit(X_train, y_train, batch_size=200, epochs=30, validation_data=(X_test, y_test), verbose=0)
        print(f'lr {lr} and momentum {mom}')
        print("-----")
        _, train_data_accuracy = model2.evaluate(X_train, y_train)
        print('Accuracy Training Data: %.2f' % (train_data_accuracy*100))

        _, test_data_accuracy = model2.evaluate(X_test, y_test)
        print('Accuracy Test Data: %.2f \n\n' % (test_data_accuracy*100))
```

lr 0.01 and momentum 0.1

```
-----  
42000/42000 [=====] - 3s 67us/sample - los  
s: 0.3867 - sparse_categorical_accuracy: 0.8815  
Accuracy Training Data: 88.15  
18000/18000 [=====] - 1s 68us/sample - los  
s: 0.6068 - sparse_categorical_accuracy: 0.8193  
Accuracy Test Data: 81.93
```

lr 0.01 and momentum 0.2

```
-----  
42000/42000 [=====] - 3s 66us/sample - los  
s: 0.2235 - sparse_categorical_accuracy: 0.9349  
Accuracy Training Data: 93.49  
18000/18000 [=====] - 1s 69us/sample - los  
s: 0.5668 - sparse_categorical_accuracy: 0.8407  
Accuracy Test Data: 84.07
```

lr 0.01 and momentum 0.3

```
-----  
42000/42000 [=====] - 3s 72us/sample - los  
s: 0.1595 - sparse_categorical_accuracy: 0.95420s - loss: 0.1605 - s  
parse_categorical_accurac  
Accuracy Training Data: 95.42  
18000/18000 [=====] - 1s 72us/sample - los  
s: 0.6084 - sparse_categorical_accuracy: 0.83940s - loss: 0.5892 - s  
parse_categorical  
Accuracy Test Data: 83.94
```

lr 0.01 and momentum 0.4

```
-----  
42000/42000 [=====] - 3s 67us/sample - los  
s: 0.1139 - sparse_categorical_accuracy: 0.9668  
Accuracy Training Data: 96.68  
18000/18000 [=====] - 1s 72us/sample - los  
s: 0.6515 - sparse_categorical_accuracy: 0.8418  
Accuracy Test Data: 84.18
```

lr 0.01 and momentum 0.5

```
-----  
42000/42000 [=====] - 3s 76us/sample - los  
s: 0.0796 - sparse_categorical_accuracy: 0.9778  
Accuracy Training Data: 97.78  
18000/18000 [=====] - 1s 79us/sample - los  
s: 0.6874 - sparse_categorical_accuracy: 0.8456  
Accuracy Test Data: 84.56
```

lr 0.01 and momentum 0.6

```
-----  
42000/42000 [=====] - 3s 73us/sample - los  
s: 0.0565 - sparse_categorical_accuracy: 0.9844  
Accuracy Training Data: 98.44  
18000/18000 [=====] - 1s 71us/sample - los  
s: 0.7082 - sparse_categorical_accuracy: 0.8479  
Accuracy Test Data: 84.79
```

lr 0.01 and momentum 0.7

```
-----  
42000/42000 [=====] - 3s 72us/sample - los  
s: 0.0663 - sparse_categorical_accuracy: 0.9802  
Accuracy Training Data: 98.02  
18000/18000 [=====] - 1s 79us/sample - los  
s: 0.7679 - sparse_categorical_accuracy: 0.8401  
Accuracy Test Data: 84.01
```

lr 0.01 and momentum 0.8

```
-----  
42000/42000 [=====] - 3s 76us/sample - los  
s: 0.0524 - sparse_categorical_accuracy: 0.9853  
Accuracy Training Data: 98.53  
18000/18000 [=====] - 1s 74us/sample - los  
s: 0.7690 - sparse_categorical_accuracy: 0.8430  
Accuracy Test Data: 84.30
```

lr 0.01 and momentum 0.9

```
-----  
42000/42000 [=====] - 3s 75us/sample - los  
s: 0.0907 - sparse_categorical_accuracy: 0.9750  
Accuracy Training Data: 97.50  
18000/18000 [=====] - 1s 81us/sample - los  
s: 0.7922 - sparse_categorical_accuracy: 0.8456  
Accuracy Test Data: 84.56
```

lr 0.02 and momentum 0.1

```
-----  
42000/42000 [=====] - 3s 73us/sample - los  
s: 0.0096 - sparse_categorical_accuracy: 0.9981  
Accuracy Training Data: 99.81  
18000/18000 [=====] - 1s 67us/sample - los  
s: 0.8032 - sparse_categorical_accuracy: 0.8610  
Accuracy Test Data: 86.10
```

lr 0.02 and momentum 0.2

```
-----  
42000/42000 [=====] - 3s 75us/sample - los  
s: 0.0071 - sparse_categorical_accuracy: 0.9989  
Accuracy Training Data: 99.89  
18000/18000 [=====] - 1s 78us/sample - los  
s: 0.8487 - sparse_categorical_accuracy: 0.8581  
Accuracy Test Data: 85.81
```

lr 0.02 and momentum 0.3

```
-----  
42000/42000 [=====] - 4s 88us/sample - los  
s: 0.0070 - sparse_categorical_accuracy: 0.9988  
Accuracy Training Data: 99.88  
18000/18000 [=====] - 1s 79us/sample - los  
s: 0.8866 - sparse_categorical_accuracy: 0.8567  
Accuracy Test Data: 85.67
```


lr 0.02 and momentum 0.4

```
-----  
42000/42000 [=====] - 3s 72us/sample - los  
s: 0.0083 - sparse_categorical_accuracy: 0.9984  
Accuracy Training Data: 99.84  
18000/18000 [=====] - 2s 91us/sample - los  
s: 0.9182 - sparse_categorical_accuracy: 0.8539  
Accuracy Test Data: 85.39
```

lr 0.02 and momentum 0.5

```
-----  
42000/42000 [=====] - 3s 69us/sample - los  
s: 0.0080 - sparse_categorical_accuracy: 0.9985  
Accuracy Training Data: 99.85  
18000/18000 [=====] - 1s 65us/sample - los  
s: 0.9339 - sparse_categorical_accuracy: 0.8528  
Accuracy Test Data: 85.28
```

lr 0.02 and momentum 0.6

```
-----  
42000/42000 [=====] - 5s 110us/sample - los  
s: 0.0129 - sparse_categorical_accuracy: 0.9964  
Accuracy Training Data: 99.64  
18000/18000 [=====] - 2s 99us/sample - los  
s: 0.9496 - sparse_categorical_accuracy: 0.8481  
Accuracy Test Data: 84.81
```

lr 0.02 and momentum 0.7

```
-----  
42000/42000 [=====] - 3s 82us/sample - los  
s: 0.0187 - sparse_categorical_accuracy: 0.9945  
Accuracy Training Data: 99.45  
18000/18000 [=====] - 2s 86us/sample - los  
s: 0.9184 - sparse_categorical_accuracy: 0.8463  
Accuracy Test Data: 84.63
```

lr 0.02 and momentum 0.8

```
-----  
42000/42000 [=====] - 3s 76us/sample - los  
s: 0.0229 - sparse_categorical_accuracy: 0.99352s - lo  
Accuracy Training Data: 99.35  
18000/18000 [=====] - 1s 72us/sample - los  
s: 0.8939 - sparse_categorical_accuracy: 0.8484  
Accuracy Test Data: 84.84
```

lr 0.02 and momentum 0.9

```
-----  
42000/42000 [=====] - 3s 76us/sample - los  
s: 0.0353 - sparse_categorical_accuracy: 0.9895  
Accuracy Training Data: 98.95  
18000/18000 [=====] - 1s 75us/sample - los  
s: 0.8375 - sparse_categorical_accuracy: 0.8474  
Accuracy Test Data: 84.74
```

lr 0.03 and momentum 0.1

```
-----  
42000/42000 [=====] - 3s 73us/sample - los  
s: 0.0038 - sparse_categorical_accuracy: 0.9996  
Accuracy Training Data: 99.96  
18000/18000 [=====] - 1s 76us/sample - los  
s: 0.8788 - sparse_categorical_accuracy: 0.8617  
Accuracy Test Data: 86.17
```

lr 0.03 and momentum 0.2

```
-----  
42000/42000 [=====] - 3s 77us/sample - los  
s: 0.0033 - sparse_categorical_accuracy: 0.9996  
Accuracy Training Data: 99.96  
18000/18000 [=====] - 1s 74us/sample - los  
s: 0.9361 - sparse_categorical_accuracy: 0.86040s - loss: 0.9168 - s  
parse_categorical_accu  
Accuracy Test Data: 86.04
```

lr 0.03 and momentum 0.3

```
-----  
42000/42000 [=====] - 3s 79us/sample - los  
s: 0.0035 - sparse_categorical_accuracy: 0.9995  
Accuracy Training Data: 99.95  
18000/18000 [=====] - 1s 77us/sample - los  
s: 0.9794 - sparse_categorical_accuracy: 0.8578  
Accuracy Test Data: 85.78
```

lr 0.03 and momentum 0.4

```
-----  
42000/42000 [=====] - 3s 83us/sample - los  
s: 0.0036 - sparse_categorical_accuracy: 0.9994  
Accuracy Training Data: 99.94  
18000/18000 [=====] - 1s 80us/sample - los  
s: 1.0047 - sparse_categorical_accuracy: 0.8576  
Accuracy Test Data: 85.76
```

lr 0.03 and momentum 0.5

```
-----  
42000/42000 [=====] - 4s 88us/sample - los  
s: 0.0032 - sparse_categorical_accuracy: 0.9996  
Accuracy Training Data: 99.96  
18000/18000 [=====] - 2s 94us/sample - los  
s: 1.0123 - sparse_categorical_accuracy: 0.8575  
Accuracy Test Data: 85.75
```

lr 0.03 and momentum 0.6

```
-----  
42000/42000 [=====] - 3s 81us/sample - los  
s: 0.0045 - sparse_categorical_accuracy: 0.9992  
Accuracy Training Data: 99.92  
18000/18000 [=====] - 2s 84us/sample - los  
s: 1.0277 - sparse_categorical_accuracy: 0.8539  
Accuracy Test Data: 85.39
```

lr 0.03 and momentum 0.7

```
-----  
42000/42000 [=====] - 4s 84us/sample - los  
s: 0.0074 - sparse_categorical_accuracy: 0.9984  
Accuracy Training Data: 99.84  
18000/18000 [=====] - 1s 80us/sample - los  
s: 0.9962 - sparse_categorical_accuracy: 0.8528  
Accuracy Test Data: 85.28
```

lr 0.03 and momentum 0.8

```
-----  
42000/42000 [=====] - 3s 83us/sample - los  
s: 0.0131 - sparse_categorical_accuracy: 0.9966  
Accuracy Training Data: 99.66  
18000/18000 [=====] - 2s 87us/sample - los  
s: 0.9555 - sparse_categorical_accuracy: 0.8511  
Accuracy Test Data: 85.11
```

lr 0.03 and momentum 0.9

```
-----  
42000/42000 [=====] - 3s 81us/sample - los  
s: 0.0345 - sparse_categorical_accuracy: 0.9891  
Accuracy Training Data: 98.91  
18000/18000 [=====] - 1s 83us/sample - los  
s: 0.8472 - sparse_categorical_accuracy: 0.8464  
Accuracy Test Data: 84.64
```

lr 0.04 and momentum 0.1

```
-----  
42000/42000 [=====] - 4s 83us/sample - los  
s: 0.0028 - sparse_categorical_accuracy: 0.9997  
Accuracy Training Data: 99.97  
18000/18000 [=====] - 2s 98us/sample - los  
s: 0.9183 - sparse_categorical_accuracy: 0.8632  
Accuracy Test Data: 86.32
```

lr 0.04 and momentum 0.2

```
-----  
42000/42000 [=====] - 3s 78us/sample - los  
s: 0.0028 - sparse_categorical_accuracy: 0.9995  
Accuracy Training Data: 99.95  
18000/18000 [=====] - 1s 78us/sample - los  
s: 0.9730 - sparse_categorical_accuracy: 0.8618  
Accuracy Test Data: 86.18
```

lr 0.04 and momentum 0.3

```
-----  
42000/42000 [=====] - 3s 73us/sample - los  
s: 0.0023 - sparse_categorical_accuracy: 0.9998  
Accuracy Training Data: 99.98  
18000/18000 [=====] - 1s 75us/sample - los  
s: 1.0210 - sparse_categorical_accuracy: 0.8604  
Accuracy Test Data: 86.04
```

lr 0.04 and momentum 0.4

```
42000/42000 [=====] - 3s 76us/sample - los
s: 0.0024 - sparse_categorical_accuracy: 0.9998
Accuracy Training Data: 99.98
18000/18000 [=====] - 1s 79us/sample - los
s: 1.0413 - sparse_categorical_accuracy: 0.8592
Accuracy Test Data: 85.92
```

lr 0.04 and momentum 0.5

```
-----
42000/42000 [=====] - 3s 83us/sample - los
s: 0.0025 - sparse_categorical_accuracy: 0.9997
Accuracy Training Data: 99.97
18000/18000 [=====] - 2s 84us/sample - los
s: 1.0401 - sparse_categorical_accuracy: 0.8601
Accuracy Test Data: 86.01
```

lr 0.04 and momentum 0.6

```
-----
42000/42000 [=====] - 4s 106us/sample - los
s: 0.0028 - sparse_categorical_accuracy: 0.9996
Accuracy Training Data: 99.96
18000/18000 [=====] - 2s 102us/sample - los
s: 1.0657 - sparse_categorical_accuracy: 0.8563
Accuracy Test Data: 85.63
```

lr 0.04 and momentum 0.7

```
-----
42000/42000 [=====] - 4s 88us/sample - los
s: 0.0048 - sparse_categorical_accuracy: 0.99882s -
Accuracy Training Data: 99.88
18000/18000 [=====] - 2s 88us/sample - los
s: 1.0556 - sparse_categorical_accuracy: 0.8556
Accuracy Test Data: 85.56
```

lr 0.04 and momentum 0.8

```
-----
42000/42000 [=====] - 4s 88us/sample - los
s: 0.0080 - sparse_categorical_accuracy: 0.9985
Accuracy Training Data: 99.85
18000/18000 [=====] - 1s 72us/sample - los
s: 0.9877 - sparse_categorical_accuracy: 0.8533
Accuracy Test Data: 85.33
```

lr 0.04 and momentum 0.9

```
-----
42000/42000 [=====] - 3s 75us/sample - los
s: 0.0303 - sparse_categorical_accuracy: 0.9912
Accuracy Training Data: 99.12
18000/18000 [=====] - 1s 80us/sample - los
s: 0.9279 - sparse_categorical_accuracy: 0.8479
Accuracy Test Data: 84.79
```

lr 0.05 and momentum 0.1

```
-----
42000/42000 [=====] - 4s 86us/sample - los
```

```
s: 0.0024 - sparse_categorical_accuracy: 0.9998
Accuracy Training Data: 99.98
18000/18000 [=====] - 1s 72us/sample - los
s: 0.9554 - sparse_categorical_accuracy: 0.8619
Accuracy Test Data: 86.19
```

lr 0.05 and momentum 0.2

```
-----
42000/42000 [=====] - 3s 82us/sample - los
s: 0.0022 - sparse_categorical_accuracy: 0.9997
Accuracy Training Data: 99.97
18000/18000 [=====] - 1s 76us/sample - los
s: 1.0107 - sparse_categorical_accuracy: 0.8621
Accuracy Test Data: 86.21
```

lr 0.05 and momentum 0.3

```
-----
42000/42000 [=====] - 3s 72us/sample - los
s: 0.0022 - sparse_categorical_accuracy: 0.9997
Accuracy Training Data: 99.97
18000/18000 [=====] - 1s 75us/sample - los
s: 1.0515 - sparse_categorical_accuracy: 0.8596
Accuracy Test Data: 85.96
```

lr 0.05 and momentum 0.4

```
-----
42000/42000 [=====] - 4s 101us/sample - los
s: 0.0020 - sparse_categorical_accuracy: 0.9998
Accuracy Training Data: 99.98
18000/18000 [=====] - 2s 98us/sample - los
s: 1.0770 - sparse_categorical_accuracy: 0.8623
Accuracy Test Data: 86.23
```

lr 0.05 and momentum 0.5

```
-----
42000/42000 [=====] - 3s 76us/sample - los
s: 0.0024 - sparse_categorical_accuracy: 0.9997
Accuracy Training Data: 99.97
18000/18000 [=====] - 1s 72us/sample - los
s: 1.0869 - sparse_categorical_accuracy: 0.8596
Accuracy Test Data: 85.96
```

lr 0.05 and momentum 0.6

```
-----
42000/42000 [=====] - 4s 93us/sample - los
s: 0.0024 - sparse_categorical_accuracy: 0.9997
Accuracy Training Data: 99.97
18000/18000 [=====] - 2s 87us/sample - los
s: 1.0896 - sparse_categorical_accuracy: 0.8612
Accuracy Test Data: 86.12
```

lr 0.05 and momentum 0.7

```
-----
42000/42000 [=====] - 4s 98us/sample - los
s: 0.0039 - sparse_categorical_accuracy: 0.9994
```

Accuracy Training Data: 99.94

18000/18000 [=====] - 2s 89us/sample - loss: 1.0959 - sparse_categorical_accuracy: 0.8562

Accuracy Test Data: 85.62

lr 0.05 and momentum 0.8

42000/42000 [=====] - 3s 81us/sample - loss: 0.0077 - sparse_categorical_accuracy: 0.9981

Accuracy Training Data: 99.81

18000/18000 [=====] - 1s 67us/sample - loss: 1.0418 - sparse_categorical_accuracy: 0.8544

Accuracy Test Data: 85.44

lr 0.05 and momentum 0.9

42000/42000 [=====] - 3s 75us/sample - loss: 0.0168 - sparse_categorical_accuracy: 0.9952

Accuracy Training Data: 99.52

18000/18000 [=====] - 1s 74us/sample - loss: 0.9097 - sparse_categorical_accuracy: 0.8547

loss: 0.8338 - sparse_categorical_accuracy: 0.8547

Accuracy Test Data: 85.47

From the statistics printed above it appears that the learning rate 0.01 and momentum 0.2 has the highest accuracy for the least loss.

We will create model with learning rate .5 and momentum .5 and see the performance of the model.

In [89]:

```
sgd = optimizers.SGD(lr=.5, momentum=.5, nesterov=False)
model2.compile(loss=losses.sparse_categorical_crossentropy, optimizer=sgd, metrics=[ 'sparse_categorical_accuracy' ])

model2.fit(X_train, y_train, batch_size=200, epochs=30, validation_data=(X_test, y_test), verbose=0)
print(f'lr .5 and momentum .5')
print("-----")
_, train_data_accuracy = model2.evaluate(X_train, y_train)
print('Accuracy Training Data: %.2f' % (train_data_accuracy*100))

_, test_data_accuracy = model2.evaluate(X_test, y_test)
print('Accuracy Test Data: %.2f \n\n' % (test_data_accuracy*100))
```

```
lr .5 and momentum .5
-----
42000/42000 [=====] - 3s 82us/sample - los
s: 0.0479 - sparse_categorical_accuracy: 0.9865
Accuracy Training Data: 98.65
18000/18000 [=====] - 1s 80us/sample - los
s: 0.7696 - sparse_categorical_accuracy: 0.8731
Accuracy Test Data: 87.31
```

When the model is build with sgd optimizer with learning rate .5 and momentum .5, the accuracy seems to have improved considerable. But the loss has increased too.

In [90]:

```
sgd = optimizers.SGD(lr=.1, momentum=.9, nesterov=False)
model2.compile(loss=losses.sparse_categorical_crossentropy, optimizer=sgd, metrics=[ 'sparse_categorical_accuracy' ])

model2.fit(X_train, y_train, batch_size=200, epochs=30, validation_data=(X_test, y_test), verbose=0)

_, train_data_accuracy = model2.evaluate(X_train, y_train)
print('Accuracy Training Data: %.2f' % (train_data_accuracy*100))

_, test_data_accuracy = model2.evaluate(X_test, y_test)
print('Accuracy Test Data: %.2f \n\n' % (test_data_accuracy*100))
```

```
42000/42000 [=====] - 3s 74us/sample - los
s: 0.0319 - sparse_categorical_accuracy: 0.9914
Accuracy Training Data: 99.14
18000/18000 [=====] - 2s 92us/sample - los
s: 0.7874 - sparse_categorical_accuracy: 0.8780
Accuracy Test Data: 87.80
```

When the model is build with sgd optimizer with learning rate .1 and momentum .9, the accuracy seems to have improved so much with the training data that it could be over fit. And the loss is significantly high.

In [116]:


```
sgd = optimizers.SGD(lr=.01, momentum=.2, nesterov=False)
model2.compile(loss=losses.sparse_categorical_crossentropy, optimizer=sgd, metrics=[ 'sparse_categorical_accuracy' ])

model2.fit(X_train, y_train, batch_size=200, epochs=30, validation_data=(X_test, y_test), verbose=0)

print(X_train.shape)
nnmodelPredictions = model2.predict_classes(X_train)

plt.figure(figsize=(10, 1))
for i in range(10):
    plt.subplot(1, 10, i+1)
    plt.imshow(X_train_original[i], cmap="gray")
    plt.axis('off')
plt.show()
print('label for each of the above image: %s' % (y_train[0:10]))
print('Model Prediction for each of the above image: %s' % nnmodelPredictions[0:10])
```

(42000, 1024)



label for each of the above image: [2 6 7 4 4 0 3 0 7 3]

Model Prediction for each of the above image: [2 6 7 4 4 0 3 0 7 3]

From this we can infer that the prediction of this model is better and more accurate. The time taken is significantly lower to build NN model as compared to the traditional classification with KNN. The accuracy of the model also significantly improves when using NN. We will store this model with the hyperparameters $lr=0.1$ and $momentum=.2$ for OCR of SVHN database.

In [98]:

```
model2.save('./svhn_ocr_model.hd5')
```

Let's load the model and predict it with a sample image to see if the OCR is as expected.

In [132]:


```
from tensorflow.keras.models import load_model
from tensorflow.keras.preprocessing import image
import cv2

model = load_model('svhn_ocr_model.hd5')
```


In [140]:

```
# Get test image ready
test_image = image.load_img('test.png', target_size=(32, 32))
plt.matshow(test_image)
test_image = image.img_to_array(test_image)
test_image = cv2.cvtColor(test_image, cv2.COLOR_BGR2GRAY)
test_image = test_image.reshape(-1, 1024) / 255
result = model.predict_classes(test_image, batch_size=1)
print(result)
```


[2]



In [141]:

```
# Get test image ready
test_image = image.load_img('test2.png', target_size=(32, 32))
plt.matshow(test_image)
test_image = image.img_to_array(test_image)
test_image = cv2.cvtColor(test_image, cv2.COLOR_BGR2GRAY)
test_image = test_image.reshape(-1, 1024) / 255
result = model.predict_classes(test_image, batch_size=1)
print(result)
```

[7]



We can conclude that the model works as expected as it seems to predicts the two random images from correctly.